

Package ‘BayesNSGP’

April 5, 2019

Title Bayesian Analysis of Non-Stationary Gaussian Process Models

Description MARK: UP FOR FILLING THIS IN? ONE SHORT PARAGRAPH.

Version 0.1.0

Date 2019-01-20

Maintainer Daniel Turek <dbt1@williams.edu>

Author Mark Risser, Daniel Turek

Depends R (>= 3.4.0)

Imports FNN,Matrix,methods,nimble,StatMatch

License GPL-3

Encoding UTF-8

LazyData true

RoxygenNote 6.1.0

R topics documented:

calcQF	2
calculateAD_ns	2
calculateU_ns	3
conditionLatentObs	4
determineNeighbors	5
dmnorm_nngp	5
dmnorm_sgv	6
inverseEigen	7
matern_corr	7
nsCorr	8
nsCrosscorr	9
nsCrossdist	10
nsCrossdist3d	11
nsDist	11
nsDist3d	12
NSGP-class	13
nsgpModel	13
orderCoordinatesMMD	14
sgvSetup	15
Index	16

calcQF	<i>Calculate the Gaussian quadratic form for the NNGP approximation</i>
--------	---

Description

calcQF calculates the quadratic form in the multivariate Gaussian based on the NNGP approximation, for a specific parameter combination. The quadratic form is $t(u)C^{-1}v$.

Usage

```
calcQF(u, v, AD, nID)
```

Arguments

u	Vector; left product.
v	Vector; right product
AD	N x (k+1) matrix; the first k columns are the 'A' matrix, and the last column is the 'D' vector. Represents the Cholesky of C^{-1} .
nID	N x k matrix of neighbor indices.

Value

A list with two components: (1) an N x 2 array containing the same spatial coordinates, ordered by MMD, and (2) the same thing, but with any NA values removed.

Examples

```
# TODO
```

calculateAD_ns	<i>Calculate A and D matrices for the NNGP approximation</i>
----------------	--

Description

calculateAD_ns calculates A and D matrices (the Cholesky of the precision matrix) needed for the NNGP approximation.

Usage

```
calculateAD_ns(dist1_3d, dist2_3d, dist12_3d, Sigma11, Sigma22, Sigma12,
  log_sigma_vec, log_tau_vec, nID, N, k, nu, d)
```

Arguments

dist1_3d	N x (k+1) x (k+1) array of distances in the x-coordinate direction.
dist2_3d	N x (k+1) x (k+1) array of distances in the y-coordinate direction.
dist12_3d	N x (k+1) x (k+1) array of cross-distances.
Sigma11	N-vector; 1-1 element of the Sigma() process.
Sigma22	N-vector; 2-2 element of the Sigma() process.
Sigma12	N-vector; 1-2 element of the Sigma() process.
log_sigma_vec	N-vector; process standard deviation values.
log_tau_vec	N-vector; nugget standard deviation values.
nID	N x k matrix of neighbor indices.
N	Scalar; number of data measurements.
k	Scalar; number of nearest neighbors.
nu	Scalar; Matern smoothness parameter.
d	TODO

Value

A N x (k+1) matrix; the first k columns are the 'A' matrix, and the last column is the 'D' vector.

Examples

```
# TODO
```

calculateU_ns	<i>Calculate the (sparse) matrix U</i>
---------------	--

Description

calculateU_ns calculates the (sparse) matrix U (i.e., the Cholesky of the inverse covariance matrix) using a nonstationary covariance function. The output only contains non-zero values and is stored as three vectors: (1) the row indices, (2) the column indices, and (3) the non-zero values. NOTE: this code assumes the all inputs correspond to the ORDERED locations.

Usage

```
calculateU_ns(dist1_3d, dist2_3d, dist12_3d, Sigma11, Sigma22, Sigma12,
  log_sigma_vec, log_tau_vec, nu, nID, cond_on_y, N, k, d, M = 0)
```

Arguments

dist1_3d	N x (k+1) x (k+1) array of distances in the x-coordinate direction.
dist2_3d	N x (k+1) x (k+1) array of distances in the y-coordinate direction.
dist12_3d	N x (k+1) x (k+1) array of cross-distances.
Sigma11	N-vector; 1-1 element of the Sigma() process.
Sigma22	N-vector; 2-2 element of the Sigma() process.

Sigma12	N-vector; 1-2 element of the Sigma() process.
log_sigma_vec	N-vector; process standard deviation values.
log_tau_vec	N-vector; nugget standard deviation values.
nu	Scalar; Matern smoothness parameter.
nID	N x k matrix of (ordered) neighbor indices.
cond_on_y	A matrix indicating whether the conditioning set for each (ordered) location is on the latent process (y, 1) or the observed values (z, 0). Calculated in sgvSetup.
N	Scalar; number of data measurements.
k	Scalar; number of nearest neighbors.
d	TODO
M	TODO

Value

TODO

Examples

```
# TODO
```

conditionLatentObs	<i>Assign conditioning sets for the SGV approximation</i>
--------------------	---

Description

conditionLatentObs assigns $q_y(i)$ vs $q_z(i)$ following Section 5.1 in Katzfuss and Guinness (2018). This function only needs to be run once per SGV analysis.

Usage

```
conditionLatentObs(nID, locs_ord, N)
```

Arguments

nID	N x k matrix of neighbor indices.
locs_ord	N x 2 matrix of locations.
N	Scalar; number of locations (observed only!).

Value

A matrix indicating whether the conditioning set for each location is on the latent process (y, 1) or the observed values (z, 0).

Examples

```
# TODO
```

determineNeighbors	<i>Determine the k-nearest neighbors for each spatial coordinate.</i>
--------------------	---

Description

determineNeighbors returns an $N \times k$ matrix of the nearest neighbors for spatial locations s , with the i th row giving indices of the k nearest neighbors to the i th location, which are selected from among the $1, \dots, (i-1)$ other spatial locations. The first row is -1's, since the first location has no neighbors. The $i=2$ through $i=(k+1)$ rows each necessarily contain $1:i$.

Usage

```
determineNeighbors(s, k)
```

Arguments

s	$N \times 2$ array of N 2-dimensional (x,y) spatial coordinates.
k	Scalar; number of neighbors

Value

An $N \times k$ matrix of nearest neighbor indices

Examples

```
# TODO
```

dmnorm_nngp	<i>Function for the evaluating the NNGP approximate density.</i>
-------------	--

Description

dmnorm_nngp (and rmnorm_nngp) calculate the approximate NNGP likelihood for a fixed set of parameters (i.e., A and D matrices). Finally, the distributions must be registered within nimble.

Usage

```
dmnorm_nngp(x, mean, AD, nID, N, k, log)
```

Arguments

x	N -vector of data.
$mean$	N -vector with current values of the mean
AD	$N \times (k+1)$ matrix; the first k columns are the ' A ' matrix, and the last column is the ' D ' vector.
nID	$N \times k$ matrix of neighbor indices.
N	Scalar; number of data measurements.
k	Scalar; number of nearest neighbors.
log	Scalar; should the density be on the log scale (1) or not (0).

Value

The NNGP approximate density.

Examples

```
# TODO
```

dmnorm_sgv	<i>Function for the evaluating the SGV approximate density.</i>
------------	---

Description

dmnorm_sgv (and rmnorm_sgv) calculate the approximate SGV likelihood for a fixed set of parameters (i.e., the U matrix). Finally, the distributions must be registered within nimble.

Usage

```
dmnorm_sgv(x, mean, U, N, k, log = 1)
```

Arguments

x	TODO
mean	TODO
U	TODO
N	TODO
k	TODO
log	TODO

Value

TODO

Examples

```
# TODO
```

inverseEigen	<i>Calculate covariance elements based on eigendecomposition components</i>
--------------	---

Description

inverseEigen calculates the inverse eigendecomposition – in other words, the covariance elements based on the eigenvalues and vectors (see Paciorek and Schervish, 2006, for details on the parameterization). The function is coded as a nimbleFunction (see the nimble package) but can also be used as a regular R function.

Usage

```
inverseEigen(eigen_comp1, eigen_comp2, eigen_comp3, which_Sigma)
```

Arguments

eigen_comp1	N-vector; contains values of the log of the second anisotropy eigenvalue for a set of locations.
eigen_comp2	N-vector; contains values of the first eigenvector component for a set of locations.
eigen_comp3	N-vector; contains values of the second eigenvector component for a set of locations.
which_Sigma	Scalar; one of (1, 2, 3), corresponding to which covariance component should be calculated (Sigma11, Sigma22, or Sigma12, respectively).

Value

A correlation matrix for a fixed set of stations and fixed parameter values.

Examples

```
# TODO
```

matern_corr	<i>Calculate a stationary Matern correlation matrix</i>
-------------	---

Description

matern_corr calculates a stationary Matern correlation matrix for a fixed set of locations, based on a range and smoothness parameter. This function is primarily used for the "npGP" and "approxGP" models. The function is coded as a nimbleFunction (see the nimble package) but can also be used as a regular R function.

Usage

```
matern_corr(dist, rho, nu)
```

Arguments

dist	N x N matrix; contains values of pairwise Euclidean distances in the x-y plane.
rho	Scalar; "range" parameter used to rescale distances
nu	Scalar; Matern smoothness parameter. $\nu = 0.5$ corresponds to the Exponential correlation; $\nu = \text{Inf}$ corresponds to the Gaussian correlation function.

Value

A correlation matrix for a fixed set of stations and fixed parameter values.

Examples

```
# TODO
```

nsCorr	<i>Calculate a nonstationary Matern correlation matrix</i>
--------	--

Description

nsCorr calculates a nonstationary correlation matrix for a fixed set of locations, based on vectors of the unique anisotropy parameters for each station. Since the correlation function uses a spatially-varying Mahalanobis distance, this function requires coordinate-specific distance matrices (see below). The function is coded as a nimbleFunction (see the nimble package) but can also be used as a regular R function.

Usage

```
nsCorr(dist1_sq, dist2_sq, dist12, Sigma11, Sigma22, Sigma12, nu, d)
```

Arguments

dist1_sq	N x N matrix; contains values of pairwise squared distances in the x-coordinate.
dist2_sq	N x N matrix; contains values of pairwise squared distances in the y-coordinate.
dist12	N x N matrix; contains values of pairwise signed cross-distances between the x- and y-coordinates. The sign of each element is important; see nsDist function for the details of this calculation. in the x-coordinate.
Sigma11	Vector of length N; contains the 1-1 element of the anisotropy process for each station.
Sigma22	Vector of length N; contains the 2-2 element of the anisotropy process for each station.
Sigma12	Vector of length N; contains the 1-2 element of the anisotropy process for each station.
nu	Scalar; Matern smoothness parameter. $\nu = 0.5$ corresponds to the Exponential correlation; $\nu = \text{Inf}$ corresponds to the Gaussian correlation function.
d	TODO

Value

A correlation matrix for a fixed set of stations and fixed parameter values.

Examples

```
# TODO
```

nsCrosscorr	<i>Calculate a nonstationary Matern cross-correlation matrix</i>
-------------	--

Description

nsCrosscorr calculates a nonstationary cross-correlation matrix between two fixed sets of locations (a prediction set with M locations, and the observed set with N locations), based on vectors of the unique anisotropy parameters for each station. Since the correlation function uses a spatially-varying Mahalanobis distance, this function requires coordinate- specific distance matrices (see below). The function is coded as a nimbleFunction (see the nimble package) but can also be used as a regular R function.

Usage

```
nsCrosscorr(Xdist1_sq, Xdist2_sq, Xdist12, Sigma11, Sigma22, Sigma12,
            PSigma11, PSigma22, PSigma12, nu, d)
```

Arguments

Xdist1_sq	M x N matrix; contains values of pairwise squared cross-distances in the x-coordinate.
Xdist2_sq	M x N matrix; contains values of pairwise squared cross-distances in the y-coordinate.
Xdist12	M x N matrix; contains values of pairwise signed cross/cross- distances between the x- and y-coordinates. The sign of each element is important; see nsDist function for the details of this calculation. in the x-coordinate.
Sigma11	Vector of length N; contains the 1-1 element of the anisotropy process for each observed location.
Sigma22	Vector of length N; contains the 2-2 element of the anisotropy process for each observed location.
Sigma12	Vector of length N; contains the 1-2 element of the anisotropy process for each observed location.
PSigma11	Vector of length N; contains the 1-1 element of the anisotropy process for each prediction location.
PSigma22	Vector of length N; contains the 2-2 element of the anisotropy process for each prediction location.
PSigma12	Vector of length N; contains the 1-2 element of the anisotropy process for each prediction location.
nu	Scalar; Matern smoothness parameter. nu = 0.5 corresponds to the Exponential correlation; nu = Inf corresponds to the Gaussian correlation function.
d	TODO

Value

A cross-correlation matrix for two fixed sets of stations and fixed parameter values.

Examples

```
# TODO
```

nsCrossdist	<i>Calculate coordinate-specific cross-distance matrices</i>
-------------	--

Description

nsCrossdist calculates coordinate-specific cross distances in x, y, and x-y for use in the nonstationary cross-correlation calculation. This function is useful for calculating posterior predictions.

Usage

```
nsCrossdist(coords, Pcoords, scale_factor = NULL, isotropic = FALSE)
```

Arguments

coords	N x 2 matrix; contains x-y coordinates of station (observed) locations.
Pcoords	M x 2 matrix; contains x-y coordinates of prediction locations.
scale_factor	Scalar; optional argument for re-scaling the distances.
isotropic	TODO

Value

A list of distances matrices, with the following components:

dist1_sq	M x N matrix; contains values of pairwise squared cross- distances in the x-coordinate.
dist2_sq	M x N matrix; contains values of pairwise squared cross- distances in the y-coordinate.
dist12	M x N matrix; contains values of pairwise signed cross- distances between the x- and y-coordinates.
scale_factor	Value of the scale factor used to rescale distances.

Examples

```
# TODO
```

nsCrossdist3d	<i>Calculate coordinate-specific distance matrices, only for nearest neighbors and store in an array</i>
---------------	--

Description

nsCrossdist3d generates and returns new 3-dimensional arrays containing the former dist1_sq, dist2_s1, and dist12 matrices, but only as needed for the k nearest-neighbors of each location. these 3D matrices (dist1_3d, dist2_3d, and dist12_3d) are used in the new implementation of calculateAD_ns().

Usage

```
nsCrossdist3d(coords, predCoords, P_nID, scale_factor = NULL,
              isotropic = FALSE)
```

Arguments

coords	N x 2 matrix; contains the x-y coordinates of stations.
predCoords	TODO
P_nID	N x k matrix; contains indices of nearest neighbors.
scale_factor	Scalar; optional argument for re-scaling the distances.
isotropic	Logical; indicates whether distances should be calculated separately for each coordinate dimension (FALSE) or simultaneously for all coordinate dimensions (TRUE). isotropic = TRUE can only be used for two-dimensional coordinate systems.

Value

Arrays with nearest neighbor distances in each coordinate direction.

Examples

```
# TODO
```

nsDist	<i>Calculate coordinate-specific distance matrices</i>
--------	--

Description

nsDist calculates x, y, and x-y distances for use in the nonstationary correlation calculation. The sign of the cross-distance is important. The function contains an optional argument for re-scaling the distances such that the coordinates lie in a square.

Usage

```
nsDist(coords, scale_factor = NULL, isotropic = FALSE)
```

Arguments

coords	N x 2 matrix; contains the x-y coordinates of stations
scale_factor	Scalar; optional argument for re-scaling the distances.
isotropic	Logical; indicates whether distances should be calculated separately for each coordinate dimension (FALSE) or simultaneously for all coordinate dimensions (TRUE). <code>isotropic = TRUE</code> can only be used for two-dimensional coordinate systems.

Value

A list of distances matrices, with the following components:

dist1_sq	N x N matrix; contains values of pairwise squared distances in the x-coordinate.
dist2_sq	N x N matrix; contains values of pairwise squared distances in the y-coordinate.
dist12	N x N matrix; contains values of pairwise signed cross- distances between the x- and y-coordinates.
scale_factor	Value of the scale factor used to rescale distances.

Examples

```
# TODO
```

nsDist3d	<i>Calculate coordinate-specific distance matrices, only for nearest neighbors and store in an array</i>
----------	--

Description

nsDist3d generates and returns new 3-dimensional arrays containing the former dist1_sq, dist2_sq, and dist12 matrices, but only as needed for the k nearest-neighbors of each location. these 3D matrices (dist1_3d, dist2_3d, and dist12_3d) are used in the new implementation of calculateAD_ns().

Usage

```
nsDist3d(coords, nID, scale_factor = NULL, isotropic = FALSE)
```

Arguments

coords	N x 2 matrix; contains the x-y coordinates of stations.
nID	N x k matrix; contains indices of nearest neighbors.
scale_factor	Scalar; optional argument for re-scaling the distances.
isotropic	Logical; indicates whether distances should be calculated separately for each coordinate dimension (FALSE) or simultaneously for all coordinate dimensions (TRUE). <code>isotropic = TRUE</code> can only be used for two-dimensional coordinate systems.

Value

Arrays with nearest neighbor distances in each coordinate direction.

Examples

```
# TODO
```

NSGP-class	<i>NSGP class</i>
------------	-------------------

Description

TODO: more detailed description here

Fields

A A is a field

nsgpModel	<i>NIMBLE code for a generic nonstationary GP model</i>
-----------	---

Description

TODO: add documentation

Usage

```
nsgpModel(tau_model = "constant", sigma_model = "constant",
  Sigma_model = "constant", mu_model = "constant",
  likelihood = "fullGP", returnModelComponents = FALSE,
  constants = list(), z, ...)
```

Arguments

tau_model	Character; specifies the model to be used for the log(tau) process. Options are "logLinReg" (log-linear regression), "mixComp" (mixture component representation), "GP" (stationary Gaussian process), and "approxGP" (approximation to a Gaussian process).
sigma_model	Character; specifies the model to be used for the log(sigma) process. See tau_model for options.
Sigma_model	Character; specifies the model to be used for the Sigma anisotropy process. Options are "covReg" (covariance regression), "compReg" (componentwise regression), "npMixComp" (nonparametric regression via the mixture component approach), "npGP" (nonparametric regression via a stationary Gaussian process), or "npApproxGP" (nonparametric regression via an approximation to a stationary Gaussian process).
mu_model	TODO
likelihood	TODO
returnModelComponents	TODO
constants	TODO
z	TODO
...	TODO

Value

A nimbleCode object.

Examples

```
# TODO
```

orderCoordinatesMMD	<i>Order coordinates according to a maximum-minimum distance criterion.</i>
---------------------	---

Description

orderCoordinatesMMD orders an array of (x,y) spatial coordinates according to the "maximum minimum distance" (MMD), as described in Guinness, 2018. (Points are selected to maximize their minimum distance to already- selected points).

Usage

```
orderCoordinatesMMD(s, exact = FALSE)
```

Arguments

s	N x 2 array of N 2-dimensional (x,y) spatial coordinates.
exact	Logical; FALSE uses a fast approximation to MMD ordering (and is almost always recommended), while TRUE uses exact MMD ordering but is infeasible for large number of locations.

Value

A list with two components: (1) an N x 2 array containing the same spatial coordinates, ordered by MMD, and (2) the same thing, but with any NA values removed.

Examples

```
# TODO
```

sgvSetup

*One-time setup wrapper function for the SGV approximation***Description**

sgvSetup is a wrapper function that sets up the SGV approximation. Three objects are required: (1) ordering the locations, (2) identify nearest neighbors, and (3) determine the conditioning set. This function only needs to be run once per SGV analysis.

Usage

```
sgvSetup(locs, locs_pred = NULL, k = 15, seed = NULL)
```

Arguments

locs	Matrix of observed locations.
locs_pred	Optional matrix of prediction locations.
k	Number of neighbors.
seed	TODO

Value

A list with the following components:

ord	A vector of ordering position for the observed locations.
ord_pred	A vector of ordering position for the prediction locations (if locs_pred is provided).
ord_all	A concatenated vector of ord and ord_pred.
locs_ord	A matrix of ordered locations (observed and prediction), included for convenience.
nID_ord	A matrix of (ordered) neighbor indices.
condition_on_y_ord	A matrix indicating whether the conditioning set for each (ordered) location is on the latent process (y, 1) or the observed values (z, 0).

Examples

```
# TODO
```

Index

`calcQF`, [2](#)
`calculateAD_ns`, [2](#)
`calculateU_ns`, [3](#)
`conditionLatentObs`, [4](#)

`determineNeighbors`, [5](#)
`dmnorm_nngp`, [5](#)
`dmnorm_sgv`, [6](#)

`inverseEigen`, [7](#)

`matern_corr`, [7](#)

`nsCorr`, [8](#)
`nsCrosscorr`, [9](#)
`nsCrossdist`, [10](#)
`nsCrossdist3d`, [11](#)
`nsDist`, [11](#)
`nsDist3d`, [12](#)
`NSGP (NSGP-class)`, [13](#)
`NSGP-class`, [13](#)
`nsgpModel`, [13](#)

`orderCoordinatesMMD`, [14](#)

`sgvSetup`, [15](#)